

<https://portal.azure.com/?>

[Microsoft Azure Education correlationId=d6682c6ed378401ba6273a093b53a511](#)

[#view/Microsoft Azure Education/JobRolesDetailBlade/jobRoleId~/5/iconUrl/https%3A%2F%2Fdocs.microsoft.com%2Fen-us%2Flearn%2Fachievements%2Fdata-ai%2Fimplement-sql-data-warehouse-trophy.svg/description/You%20design%20and%20implement%20the%20management%2C%20monitoring%2C%20and%20security%20of%20data%20solutions.%20Get%20the%20skills%20and%20knowledge%20needed%20to%20build%20your%20career%20as%20a%20successful%20Data%20Engineer./isActive~/true/isFeatured~/true/locale//name/Data%20Engineer/rank~/5/role/Data%20Engineer](#)

## The Rise of Streaming Analytics


Batch processing stopped working because:

Batch jobs were taking longer and needed more processing power

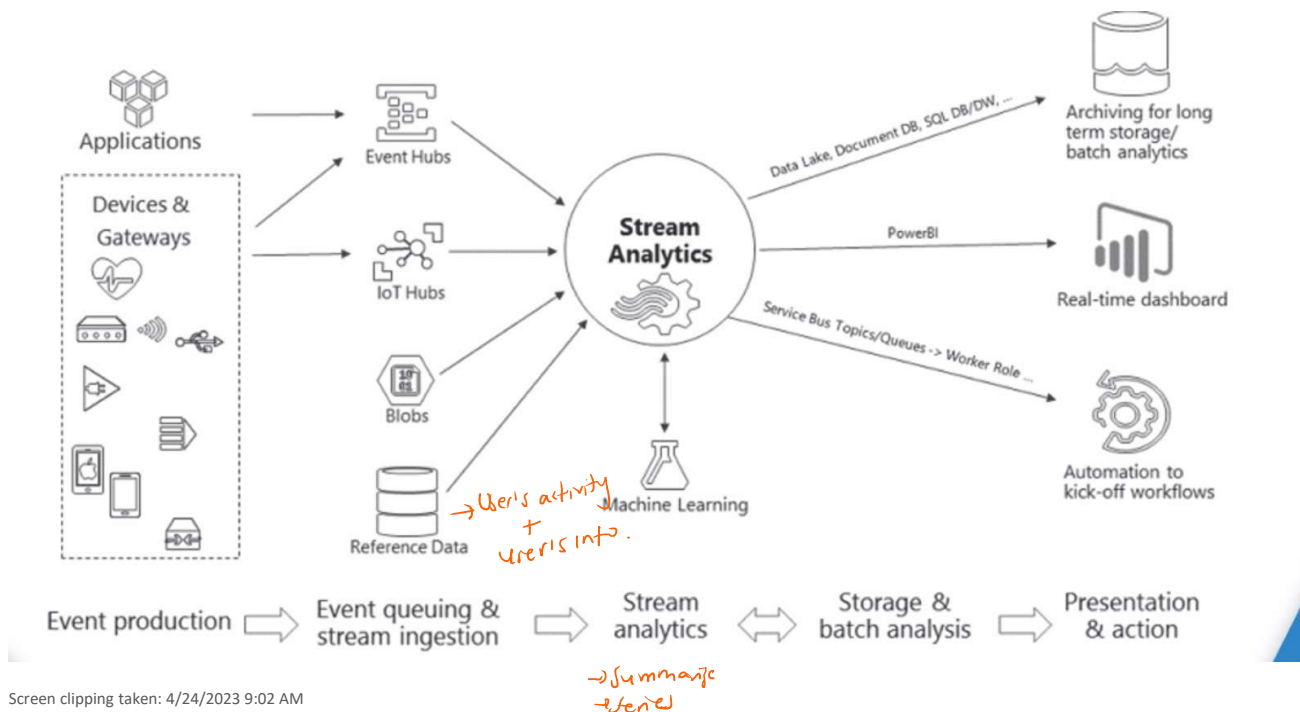
Companies had applications that needed data to be processed in real time

Screen clipping taken: 4/24/2023 8:54 AM

## Typical Uses



Screen clipping taken: 4/24/2023 8:55 AM



Screen clipping taken: 4/24/2023 9:02 AM

# Alternatives

## Spark or Storm on HDInsight

- Cluster
- Hadoop ecosystem
- Scala, Java, or Python

## Azure Streaming Analytics

- Job
- Microsoft ecosystem
- Stream Analytics Query Language (variant of T-SQL)

Screen clipping taken: 4/24/2023 9:04 AM

Question

CORRECT

Which of the following is not a required step to parallelize your Azure Stream Analytics job?

☐ Write the output to the n partitions.

☐ Write the query to act on the n partitions.

☒ Configure your job to run on exactly n nodes.

☐ Split the input into n partitions.

☐ I don't know

Explanation

The tricky part is that you need to configure your job in such a way that you can split it into multiple pieces. To achieve that, you need to do three things. First, the input must be split into n partitions. Second, the query must be written to act on those n partitions. Third, the output must be written to n partitions.

Learn more: </course/introduction-to-azure-stream-analytics/scaling-100/>

Something wrong with this question? [Report an issue](#)

Bookmark

Screen clipping taken: 4/27/2023 8:35 PM

<

Question 2

>

Question

CORRECT

Which of the following statements about logs in Azure Stream Analytics is true?

Activity logs can be enabled in the Monitoring menu.

Diagnostic logs are always available by default, without enabling them.

Diagnostic logs can be enabled in the Monitoring menu.

Diagnostic logs only include entries for high-level operations.

I don't know

Explanation

Activity logs are always available, but they only include entries for high-level operations. To get more details, you can enable diagnostic logs, which are in the Monitoring menu.

Bookmark

Learn more: </course/introduction-to-azure-stream-analytics/troubleshooting-100/>

Something wrong with this question?

Report an issue

Screen clipping taken: 4/27/2023 8:36 PM

<

Question 2

>

Question

CORRECT

Which of the following statements about logs in Azure Stream Analytics is true?

Activity logs can be enabled in the Monitoring menu.

Diagnostic logs are always available by default, without enabling them.

Diagnostic logs can be enabled in the Monitoring menu.

Diagnostic logs only include entries for high-level operations.

I don't know

Explanation

Activity logs are always available, but they only include entries for high-level operations. To get more details, you can enable diagnostic logs, which are in the Monitoring menu.

Bookmark

Learn more: </course/introduction-to-azure-stream-analytics/troubleshooting-100/>

Screen clipping taken: 4/27/2023 8:37 PM

### Question 3

Question

CORRECT

To prevent your Azure Stream Analytics job from running out of memory, you should create an alert that tells you if the SU utilization goes above \_\_\_\_%.

☐ 50

☐ 90

☒ 80

☐ 20

☐ I don't know

Explanation

If your job runs out of memory, then it will fail. To prevent this from happening, you should create an alert that tells you if the SU utilization goes above 80%.

Bookmark

Learn more: </course/introduction-to-azure-stream-analytics/scaling-100/>

Something wrong with this question? [Report an issue](#)

### Question 4

Question

CORRECT

Which type of logs in Azure Stream Analytics contain connectivity errors and data processing errors?

☐ authoring logs

☒ execution logs

☐ static logs

☐ connection logs

☐ I don't know

Explanation

There are two types of logs. Execution logs contain entries about events that happen during job execution, especially errors, such as connectivity errors and data processing errors. This is a good one to enable. Authoring logs contain entries related to job authoring operations, such as creating a job, adding inputs and outputs, adding the query, and starting and stopping the job. In addition, you can enable logging of all metrics, which could be helpful.

Bookmark

Learn more: </course/introduction-to-azure-stream-analytics/troubleshooting-100/>

Screen clipping taken: 4/27/2023 8:41 PM

Question 5

Question

 CORRECT

In Azure Stream Analytics, which part of a Streaming Unit is most important?

☒ memory

☐ input

☐ output

☐ CPU

☐ I don't know

Explanation

An SU represents a certain capacity of CPU, memory, and I/O. The most important of these is memory because Stream Analytics does all of its processing in memory rather than on disk.

 Bookmark

Learn more: </course/introduction-to-azure-stream-analytics/scaling-100/>

Something wrong with this question?  Report an issue

Screen clipping taken: 4/27/2023 8:41 PM